

Untitled

July 13, 2021

1 Project 2 : Healthcare

- DESCRIPTION

Problem Statement

- NIDDK (National Institute of Diabetes and Digestive and Kidney Diseases) research creates knowledge about and treatments for the most chronic, costly, and consequential diseases.
- The dataset used in this project is originally from NIDDK. The objective is to predict whether or not a patient has diabetes, based on certain diagnostic measurements included in the dataset.
- Build a model to accurately predict whether the patients in the dataset have diabetes or not.

Dataset Description

- The datasets consists of several medical predictor variables and one target variable (Outcome). Predictor variables includes the number of pregnancies the patient has had, their BMI, insulin level, age, and more.

[Variables||Description| |-----| |-----| |Pregnancies||Number of times pregnant **Numeric**| |Glucose||Plasma glucose concentration in an oral glucose tolerance test **Numeric**| |BloodPressure||Diastolic blood pressure (mm Hg) **Numeric**| |SkinThickness||Triceps skinfold thickness (mm) **Numeric**| |Insulin||Two hour serum insulin **Numeric**| |BMI||Body Mass Index **Numeric**| |DiabetesPedigreeFunction||Diabetes pedigree function **Numeric**| |Age||Age in years **Numeric**| |Outcome||Class variable (either 0 or 1). 268 of 768 values are 1, and the others are 0|

1.0.1 Week 1 Task

- Perform descriptive analysis. Understand the variables and their corresponding values. On the columns below, a value of zero does not make sense and thus indicates missing value: 1)Glucose 2)BloodPressure 3)SkinThickness 4)Insulin 5)BMI
- Visually explore these variables using histograms. Treat the missing values accordingly.
- There are integer and float data type variables in this dataset. Create a count (frequency) plot describing the data types and the count of variables.

```
[1]: # Importing necessary libraries
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
```

```
[2]: healthcare_data = pd.read_csv("health care diabetes.csv")
healthcare_data.head()
```

```
[2]:
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI \
0	6	148	72	35	0	33.6
1	1	85	66	29	0	26.6
2	8	183	64	0	0	23.3
3	1	89	66	23	94	28.1
4	0	137	40	35	168	43.1

	DiabetesPedigreeFunction	Age	Outcome
0	0.627	50	1
1	0.351	31	0
2	0.672	32	1
3	0.167	21	0
4	2.288	33	1

```
[3]: healthcare_data.count() # Total 768 rows are present in the dataset.
healthcare_data.shape      # (768, 9) shape of the data.
```

```
[3]: (768, 9)
```

```
[4]: healthcare_data.describe() # description of the all features in dataset.
```

```
[4]:
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin \
count	768.000000	768.000000	768.000000	768.000000	768.000000
mean	3.845052	120.894531	69.105469	20.536458	79.799479
std	3.369578	31.972618	19.355807	15.952218	115.244002
min	0.000000	0.000000	0.000000	0.000000	0.000000
25%	1.000000	99.000000	62.000000	0.000000	0.000000
50%	3.000000	117.000000	72.000000	23.000000	30.500000
75%	6.000000	140.250000	80.000000	32.000000	127.250000
max	17.000000	199.000000	122.000000	99.000000	846.000000

	BMI	DiabetesPedigreeFunction	Age	Outcome
count	768.000000	768.000000	768.000000	768.000000
mean	31.992578	0.471876	33.240885	0.348958
std	7.884160	0.331329	11.760232	0.476951
min	0.000000	0.078000	21.000000	0.000000
25%	27.300000	0.243750	24.000000	0.000000
50%	32.000000	0.372500	29.000000	0.000000

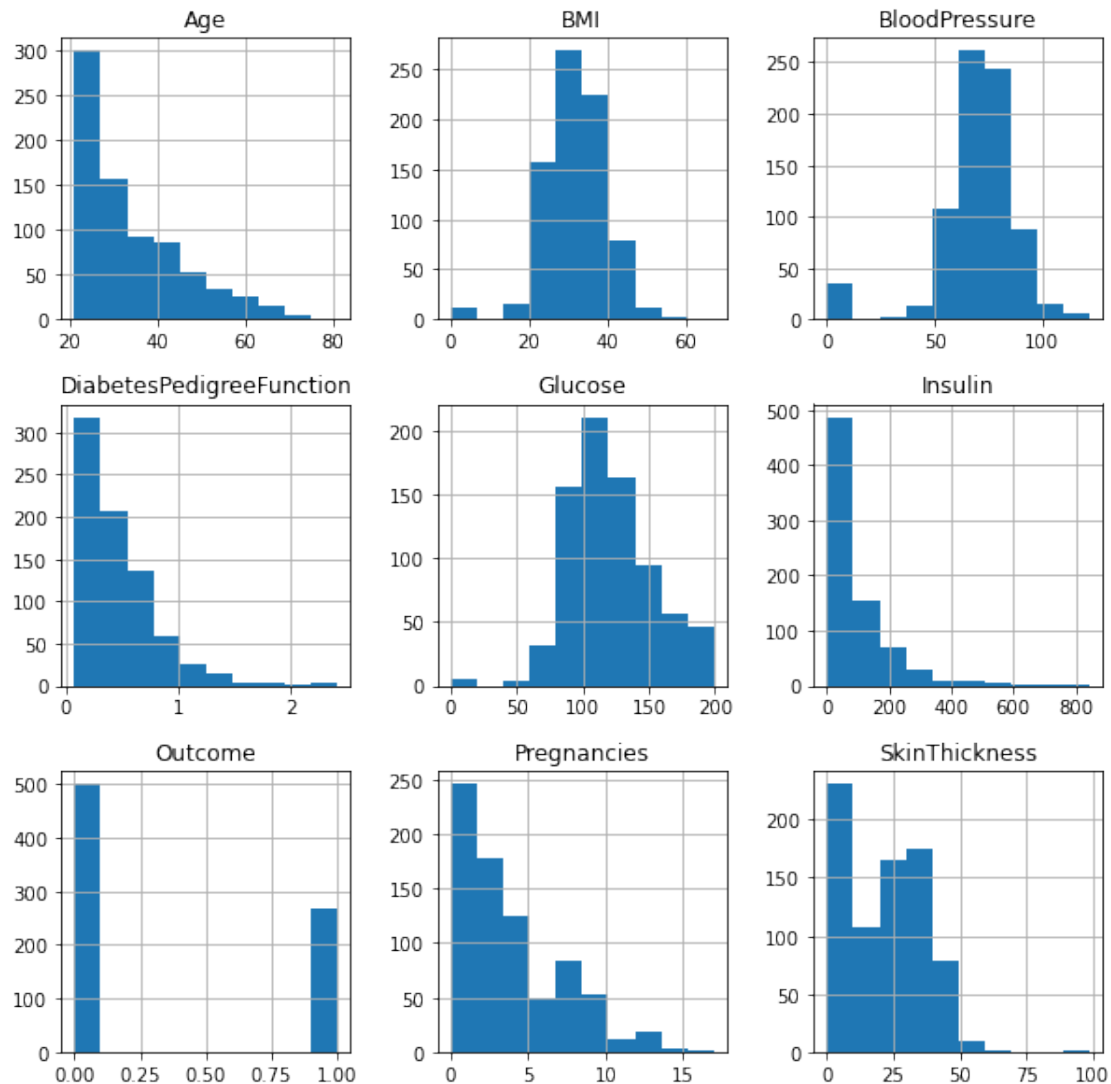
75%	36.600000	0.626250	41.000000	1.000000
max	67.100000	2.420000	81.000000	1.000000

```
[5]: healthcare_data.dtypes
```

```
[5]: Pregnancies          int64
      Glucose             int64
      BloodPressure       int64
      SkinThickness       int64
      Insulin             int64
      BMI                 float64
      DiabetesPedigreeFunction float64
      Age                int64
      Outcome            int64
      dtype: object
```

```
[6]: healthcare_data.hist(figsize=(10,10))
```

```
[6]: array([[<matplotlib.axes._subplots.AxesSubplot object at 0x000002D8F867F6A0>,
             <matplotlib.axes._subplots.AxesSubplot object at 0x000002D8F8D5AB50>,
             <matplotlib.axes._subplots.AxesSubplot object at 0x000002D8F8D89FA0>],
            [<matplotlib.axes._subplots.AxesSubplot object at 0x000002D8F8DC3460>,
             <matplotlib.axes._subplots.AxesSubplot object at 0x000002D8F8DEE8B0>,
             <matplotlib.axes._subplots.AxesSubplot object at 0x000002D8F8E19C40>],
            [<matplotlib.axes._subplots.AxesSubplot object at 0x000002D8F8E19D30>,
             <matplotlib.axes._subplots.AxesSubplot object at 0x000002D8F8E53220>,
             <matplotlib.axes._subplots.AxesSubplot object at 0x000002D8F8EADA30>]],
      dtype=object)
```



```
[7]: healthcare_data.isnull().sum()
```

```
[7]: Pregnancies      0
      Glucose          0
      BloodPressure    0
      SkinThickness    0
      Insulin          0
      BMI              0
      DiabetesPedigreeFunction  0
      Age              0
      Outcome          0
      dtype: int64
```

Note: 0 value present in Glucose,bloodpressure,skintickness,insulin,BMI doesn't make sense. We

will replace this 0 with median of the particular column.

```
[8]: # Missing values treatment
healthcare_data['Glucose']=healthcare_data['Glucose'].
    ↪replace(0,healthcare_data['Glucose'].median())
healthcare_data['BloodPressure']=healthcare_data['BloodPressure'].
    ↪replace(0,healthcare_data['BloodPressure'].median())
healthcare_data['SkinThickness']=healthcare_data['SkinThickness'].
    ↪replace(0,healthcare_data['SkinThickness'].median())
healthcare_data['Insulin']=healthcare_data['Insulin'].
    ↪replace(0,healthcare_data['Insulin'].median())
healthcare_data['BMI']=healthcare_data['BMI'].replace(0,healthcare_data['BMI'].
    ↪median())
```

```
[9]: healthcare_data.head(10) # missing value treatment done successfully.
```


```
[9]:
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	\
0	6	148	72	35	30.5	33.6	
1	1	85	66	29	30.5	26.6	
2	8	183	64	23	30.5	23.3	
3	1	89	66	23	94.0	28.1	
4	0	137	40	35	168.0	43.1	
5	5	116	74	23	30.5	25.6	
6	3	78	50	32	88.0	31.0	
7	10	115	72	23	30.5	35.3	
8	2	197	70	45	543.0	30.5	
9	8	125	96	23	30.5	32.0	

	DiabetesPedigreeFunction	Age	Outcome
0	0.627	50	1
1	0.351	31	0
2	0.672	32	1
3	0.167	21	0
4	2.288	33	1
5	0.201	30	0
6	0.248	26	1
7	0.134	29	0
8	0.158	53	1
9	0.232	54	1

```
[10]: #Create a count (frequency) plot describing the data types and the count of
    ↪variables.
plt.figure(figsize=(20,10))
sns.countplot(x='Age',data=healthcare_data)
```

```
[10]: <matplotlib.axes._subplots.AxesSubplot at 0x2d8f9087eb0>
```



1.0.2 Project Task: Week 2

Data Exploration:

1. Check the balance of the data by plotting the count of outcomes by their value. Describe your findings and plan future course of action.
2. Create scatter charts between the pair of variables to understand the relationships. Describe your findings.
3. Perform correlation analysis. Visually explore it using a heat map.

```
[11]: healthcare_data.dtypes
```


```
[11]: Pregnancies      int64
      Glucose          int64
      BloodPressure    int64
      SkinThickness    int64
      Insulin          float64
      BMI              float64
      DiabetesPedigreeFunction float64
      Age              int64
      Outcome          int64
      dtype: object
```

observation :Here the **Outcome** is numeric we will convert this into categorical.

```
[12]: healthcare_data['Outcome']=pd.Categorical(healthcare_data['Outcome'])
```

```
[13]: # Check the balance of the data by plotting the count of outcomes by their
      ↪ value. Describe your findings
      print(healthcare_data['Outcome'].value_counts())
      healthcare_data['Outcome'].value_counts().plot.bar();
```

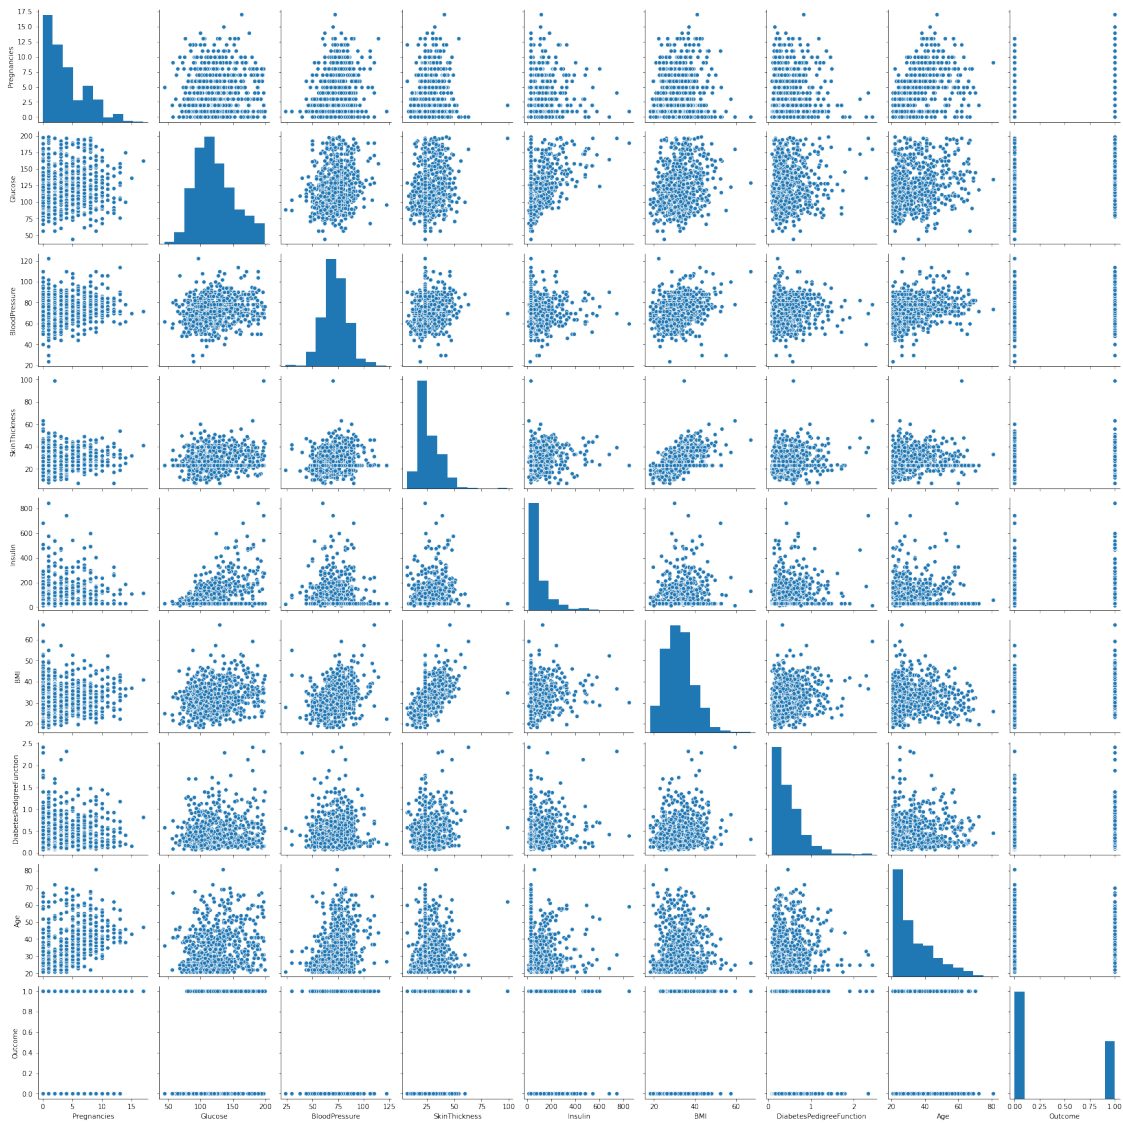
```
0    500
1    268
Name: Outcome, dtype: int64
```



Observation : The data is balanced with 500 zeros and 268 ones.

```
[14]: # Create scatter charts between the pair of variables to understand the
      ↪ relationships. Describe your findings.
      sns.pairplot(healthcare_data,palette='hue')
```

```
[14]: <seaborn.axisgrid.PairGrid at 0x2d8f9625f40>
```



```
[15]: #Perform correlation analysis. Visually explore it using a heat map.
healthcare_data.corr()
```

```
[15]:
```


	Pregnancies	Glucose	BloodPressure	SkinThickness	\
Pregnancies	1.000000	0.128213	0.208615	0.032568	
Glucose	0.128213	1.000000	0.218937	0.172143	
BloodPressure	0.208615	0.218937	1.000000	0.147809	
SkinThickness	0.032568	0.172143	0.147809	1.000000	
Insulin	-0.055697	0.357573	-0.028721	0.238188	
BMI	0.021546	0.231400	0.281132	0.546951	
DiabetesPedigreeFunction	-0.033523	0.137327	-0.002378	0.142977	
Age	0.544341	0.266909	0.324915	0.054514	

	Insulin	BMI	DiabetesPedigreeFunction	\
Pregnancies	-0.055697	0.021546		-0.033523
Glucose	0.357573	0.231400		0.137327
BloodPressure	-0.028721	0.281132		-0.002378
SkinThickness	0.238188	0.546951		0.142977
Insulin	1.000000	0.189022		0.178029
BMI	0.189022	1.000000		0.153506
DiabetesPedigreeFunction	0.178029	0.153506		1.000000
Age	-0.015413	0.025744		0.033561

	Age
Pregnancies	0.544341
Glucose	0.266909
BloodPressure	0.324915
SkinThickness	0.054514
Insulin	-0.015413
BMI	0.025744
DiabetesPedigreeFunction	0.033561
Age	1.000000

```
[16]: plt.figure(figsize=(10,8))
      sns.heatmap(healthcare_data.corr(),annot=True)
```

```
[16]: <matplotlib.axes._subplots.AxesSubplot at 0x2d8fb01d1c0>
```



Observation : There is no strong linear relation in between the feature variables.

1.0.3 Project Task: Week 3

Data Modeling:

1. Devise strategies for model building. It is important to decide the right validation framework. Express your thought process.
2. Apply an appropriate classification algorithm to build a model. Compare various models with the results from KNN algorithm.
3. Data Modeling: Create a classification report by analyzing sensitivity, specificity, AUC (ROC curve), etc. Please be descriptive to explain what values of these parameter you have used.

Devise strategies for model building. It is important to decide the right validation framework. Express your thought process.

- First we will split the data using train test split. Here validation framework is train test split because the Target column data is balanced. I prefer k-fold cross validation technique when there is highly bias data. If I use train test split in this case then there will be chances of missing the data points while training the model.
- will create the model using different classification algorithm
- compare accuracy of every model.
- will go for higher accuracy appropriate model.

```
[17]: # divide the data
X=healthcare_data.drop('Outcome',axis=1)
y=healthcare_data['Outcome']
```

```
[18]: # train test split
from sklearn.model_selection import train_test_split
X_train,X_test,y_train,y_test = train_test_split(X,y,test_size=0.
↪2,random_state=21)
```

```
[19]: print(X_train.shape)
print(X_test.shape)
print(y_train.shape)
print(y_test.shape)
```

```
(614, 8)
(154, 8)
(614,)
(154,)
```

```
[20]: #Apply an appropriate classification algorithm to build a model. Compare
↪various models with the results from KNN algorithm.
from sklearn.neighbors import KNeighborsClassifier
knn_model = KNeighborsClassifier(n_neighbors=4)
knn_model.fit(X_train,y_train)
```

```
[20]: KNeighborsClassifier(n_neighbors=4)
```

```
[21]: # prediction from training and the testing data to know model accuracy
knn_y_pred_train= knn_model.predict(X_train)
knn_y_pred_test = knn_model.predict(X_test)
```

```
[22]: # Check whether the model is overfitted,underfitted or it is appropriate
from sklearn.metrics import accuracy_score,classification_report
print('The accuracy of the training model is :
↪',accuracy_score(y_train,knn_y_pred_train))
print('The accuracy of the testing model is :
↪',accuracy_score(y_test,knn_y_pred_test))
```

```
The accuracy of the training model is : 0.8143322475570033
The accuracy of the testing model is : 0.6688311688311688
```

```
[23]: print(classification_report(y_test,knn_y_pred_test))
```

	precision	recall	f1-score	support
0	0.66	0.93	0.77	94
1	0.70	0.27	0.39	60
accuracy			0.67	154
macro avg	0.68	0.60	0.58	154
weighted avg	0.68	0.67	0.62	154

1.0.4 Using Logistic Regression

```
[24]: from sklearn.linear_model import LogisticRegression
log_model = LogisticRegression()
log_model.fit(X_train,y_train)
```

C:\Users\nisarg\anaconda3\lib\site-packages\sklearn\linear_model_logistic.py:762: ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```
n_iter_i = _check_optimize_result(
```

```
[24]: LogisticRegression()
```

```
[25]: # prediction from training and the testing data to know model accuracy
log_y_pred_train= log_model.predict(X_train)
log_y_pred_test = log_model.predict(X_test)
```

```
[26]: print('The accuracy of the training model is :
      ↪',accuracy_score(y_train,log_y_pred_train))
print('The accuracy of the testing model is :
      ↪',accuracy_score(y_test,log_y_pred_test))
```

The accuracy of the training model is : 0.7915309446254072

The accuracy of the testing model is : 0.7402597402597403

```
[27]: print(classification_report(y_test,log_y_pred_test))
```

	precision	recall	f1-score	support
0	0.73	0.90	0.81	94
1	0.76	0.48	0.59	60
accuracy			0.74	154
macro avg	0.75	0.69	0.70	154
weighted avg	0.74	0.74	0.72	154

Observation : Training accuracy of knn is higher than logistic but testing accuracy of logistic is greater than knn.

1.0.5 Using Decision Tree

```
[28]: from sklearn.tree import DecisionTreeClassifier
      dec_model = DecisionTreeClassifier()
      dec_model.fit(X_train,y_train)
```

```
[28]: DecisionTreeClassifier()
```

```
[29]: # prediction from training and the testing data to know model accuracy
      dec_y_pred_train= dec_model.predict(X_train)
      dec_y_pred_test = dec_model.predict(X_test)
```

```
[30]: print('The accuracy of the training model is :
      ↪',accuracy_score(y_train,dec_y_pred_train))
      print('The accuracy of the testing model is :
      ↪',accuracy_score(y_test,dec_y_pred_test))
```

The accuracy of the training model is : 1.0

The accuracy of the testing model is : 0.7012987012987013

```
[31]: print(classification_report(y_test,dec_y_pred_test))
```

	precision	recall	f1-score	support
0	0.74	0.79	0.76	94
1	0.63	0.57	0.60	60
accuracy			0.70	154
macro avg	0.68	0.68	0.68	154
weighted avg	0.70	0.70	0.70	154

Observation : Here both training and testing accuracy is greater than knn.

1.0.6 Using Random forest

```
[32]: from sklearn.ensemble import RandomForestClassifier
ref_model = RandomForestClassifier(n_estimators=150)
ref_model.fit(X_train,y_train)
```

```
[32]: RandomForestClassifier(n_estimators=150)
```

```
[33]: # prediction from training and the testing data to know model accuracy
ref_y_pred_train= ref_model.predict(X_train)
ref_y_pred_test = ref_model.predict(X_test)
```

```
[34]: print('The accuracy of the training model is :
      ↪',accuracy_score(y_train,ref_y_pred_train))
print('The accuracy of the testing model is :
      ↪',accuracy_score(y_test,ref_y_pred_test))
```

The accuracy of the training model is : 1.0

The accuracy of the testing model is : 0.7402597402597403

```
[35]: print(classification_report(y_test,ref_y_pred_test))
```

	precision	recall	f1-score	support
0	0.73	0.91	0.81	94
1	0.78	0.47	0.58	60
accuracy			0.74	154
macro avg	0.75	0.69	0.70	154
weighted avg	0.75	0.74	0.72	154

Observation : Here than both the accuracy is higher than knn and decision tree.

1.1 Hyperparameter tuning

```
[36]: from sklearn.model_selection import RandomizedSearchCV
Ran_ref_model = RandomForestClassifier()
```

```
[37]: n_estimators= [int(x) for x in np.linspace(100,1000,10)]
criterion=["gini", "entropy"]
max_depth=[4,5,6]
min_samples_split=[2,3,4]
min_samples_leaf=[1,2,3]
```

```
[38]: param = {'n_estimators':n_estimators,'criterion':criterion,'max_depth':
↳max_depth,'min_samples_split':min_samples_split,'min_samples_leaf':
↳min_samples_leaf}
```

```
[39]: ran_ref_final_model =_
↳RandomizedSearchCV(Ran_ref_model,param_distributions=param,n_iter=50,cv=4,n_jobs=-1,verbose
ran_ref_final_model.fit(X_train,y_train)
```

Fitting 4 folds for each of 50 candidates, totalling 200 fits

[Parallel(n_jobs=-1)]: Using backend LokyBackend with 4 concurrent workers.

[Parallel(n_jobs=-1)]: Done 42 tasks | elapsed: 18.7s

[Parallel(n_jobs=-1)]: Done 192 tasks | elapsed: 1.4min

[Parallel(n_jobs=-1)]: Done 200 out of 200 | elapsed: 1.5min finished

```
[39]: RandomizedSearchCV(cv=4, estimator=RandomForestClassifier(), n_iter=50,
n_jobs=-1,
param_distributions={'criterion': ['gini', 'entropy'],
'max_depth': [4, 5, 6],
'min_samples_leaf': [1, 2, 3],
'min_samples_split': [2, 3, 4],
'n_estimators': [100, 200, 300, 400,
500, 600, 700, 800,
900, 1000]},
verbose=1)
```

```
[40]: ran_ref_final_model.best_estimator_
```

```
[40]: RandomForestClassifier(max_depth=6, min_samples_leaf=3, min_samples_split=3,
n_estimators=500)
```

```
[41]: # prediction from training and the testing data to know model accuracy
ran_ref_y_pred_train= ran_ref_final_model.predict(X_train)
ran_ref_y_pred_test = ran_ref_final_model.predict(X_test)
```

```
[42]: print('The accuracy of the training model is :
↳',accuracy_score(y_train,ran_ref_y_pred_train))
print('The accuracy of the testing model is :
↳',accuracy_score(y_test,ran_ref_y_pred_test))
```

The accuracy of the training model is : 0.8583061889250815

The accuracy of the testing model is : 0.7402597402597403

```
[43]: print(classification_report(y_test,ran_ref_y_pred_test))
```

	precision	recall	f1-score	support
0	0.73	0.91	0.81	94

	1	0.78	0.47	0.58	60
accuracy				0.74	154
macro avg		0.75	0.69	0.70	154
weighted avg		0.75	0.74	0.72	154

Observation : Using the random Forest Classifier with hyper parameter tuning we are getting the appropriate model where accuracy of both train and test is higher.

Data Modeling: Create a classification report by analyzing sensitivity, specificity, AUC (ROC curve), etc. Please be descriptive to explain what values of these parameter you have used.

```
[44]: from sklearn.metrics import roc_curve,roc_auc_score
      # Taking prediction probability from every model to calculate fpr,tpr,threshold
      ↪values to draw roc curve
      y_knn_pred = knn_model.predict_proba(X_test)[: ,1]
      y_log_pred = log_model.predict_proba(X_test)[: ,1]
      y_dec_pred = dec_model.predict_proba(X_test)[: ,1]
      y_ranfor_pred = ran_ref_final_model.predict_proba(X_test)[: ,1]
```

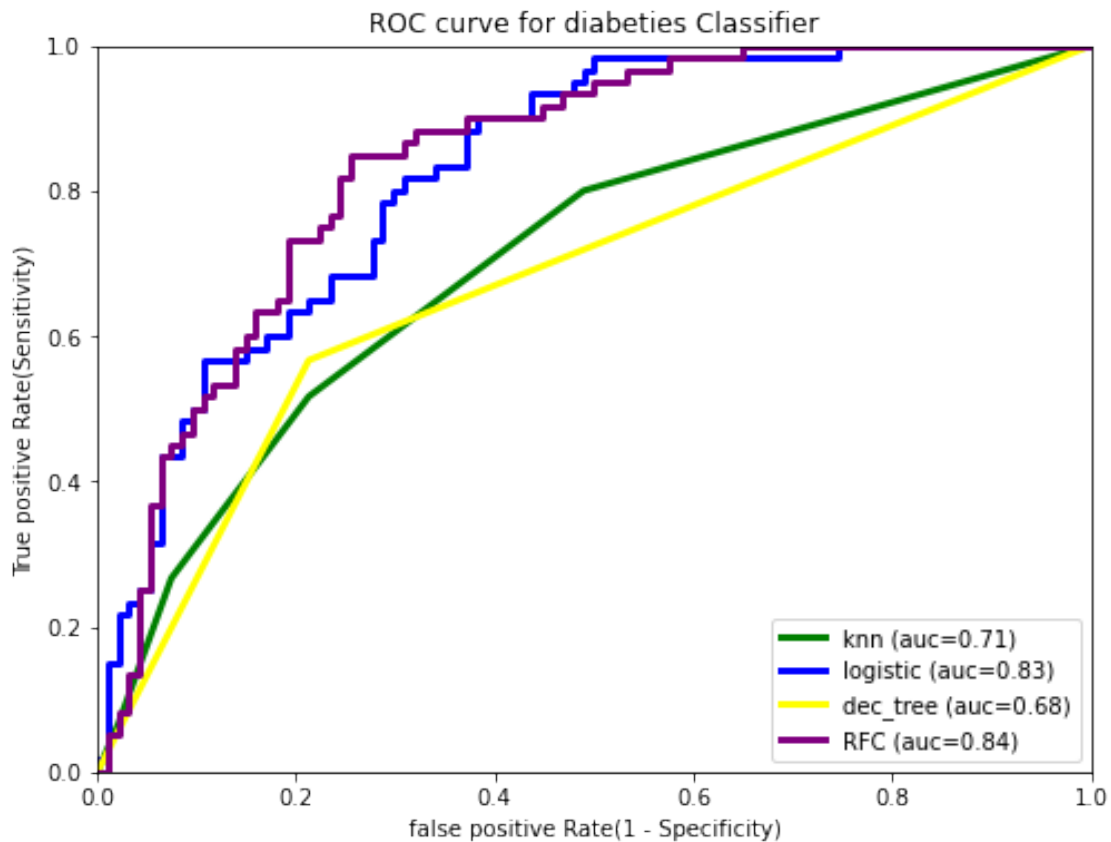
```
[45]: # calculation of fpr(False positive rate), tpr(True positive rate) values for
      ↪every model
      fpr_knn,tpr_knn,threshold_knn =
      ↪roc_curve(y_test,y_knn_pred,drop_intermediate=False)
      fpr_log,tpr_log,threshold_log =
      ↪roc_curve(y_test,y_log_pred,drop_intermediate=False)
      fpr_dec,tpr_dec,threshold_dec =
      ↪roc_curve(y_test,y_dec_pred,drop_intermediate=False)
      fpr_ran,tpr_ran,threshold_ran =
      ↪roc_curve(y_test,y_ranfor_pred,drop_intermediate=False)
```

```
[46]: # Calculating AUC values for every model
      auc_knn = roc_auc_score(y_test,y_knn_pred)
      auc_log = roc_auc_score(y_test,y_log_pred)
      auc_dec = roc_auc_score(y_test,y_dec_pred)
      auc_ran = roc_auc_score(y_test,y_ranfor_pred)
```

```
[48]: # Plotting roc curve to check which model is gives higher
      ↪accuracy,recall,precision
      plt.figure(figsize=(8,6))
      plt.xlim([0.0,1.0])
      plt.ylim([0.0,1.0])
      plt.title("ROC curve for diabeties Classifier")
      plt.xlabel('false positive Rate(1 - Specificity)')
      plt.ylabel("True positive Rate(Sensitivity)")
```



```
plt.plot(fpr_knn,tpr_knn,color='green',lw=3,label='knn (auc=%0.2f)' %auc_knn)
plt.plot(fpr_log,tpr_log,color='blue',lw=3,label='logistic (auc=%0.2f)' %auc_log)
plt.plot(fpr_dec,tpr_dec,color='yellow',lw=3,label='dec_tree (auc=%0.2f)' %
→%auc_dec)
plt.plot(fpr_ran,tpr_ran,color='purple',lw=3,label='RFC (auc=%0.2f)' %auc_ran)
plt.legend()
plt.show()
```



Observation : from the above roc_auc curve we can say that for the above problem Logistic regression and random forest model gives us better output with auc score of 0.83 and 0.84 respectively

1.1.1 Project Task: Week 4

Data Reporting:

2. Create a dashboard in tableau by choosing appropriate chart types and metrics useful for the business. The dashboard must entail the following:
 - a. Pie chart to describe the diabetic or non-diabetic population

- b. Scatter charts between relevant variables to analyze the relationships
- c. Histogram or frequency charts to analyze the distribution of the data
- d. Heatmap of correlation analysis among the relevant variables
- e. Create bins of these age values: 20-25, 25-30, 30-35, etc. Analyze different variables for these age brackets using a bubble chart.

```
[41]: healthcare_data.to_csv(r'C:\Users\nisarg\Desktop\Data Science\Capstone_1  
↳Projects\Project 2\Healthcare - Diabetes\tableaufile.csv',index=False)
```

Observation : we have converted the file to csv for visualization purpose

Analysis of Diabetic Patients

Pie chart to describe the diabetic or non-di..	Scatter charts between relevant var..	Histogram plot to see the distribution of th..	Operation on age column by creating th..
--	---------------------------------------	--	--

Diabeties

■


 Diabetic

■

 NonDiabetic


Count of Outcome

768



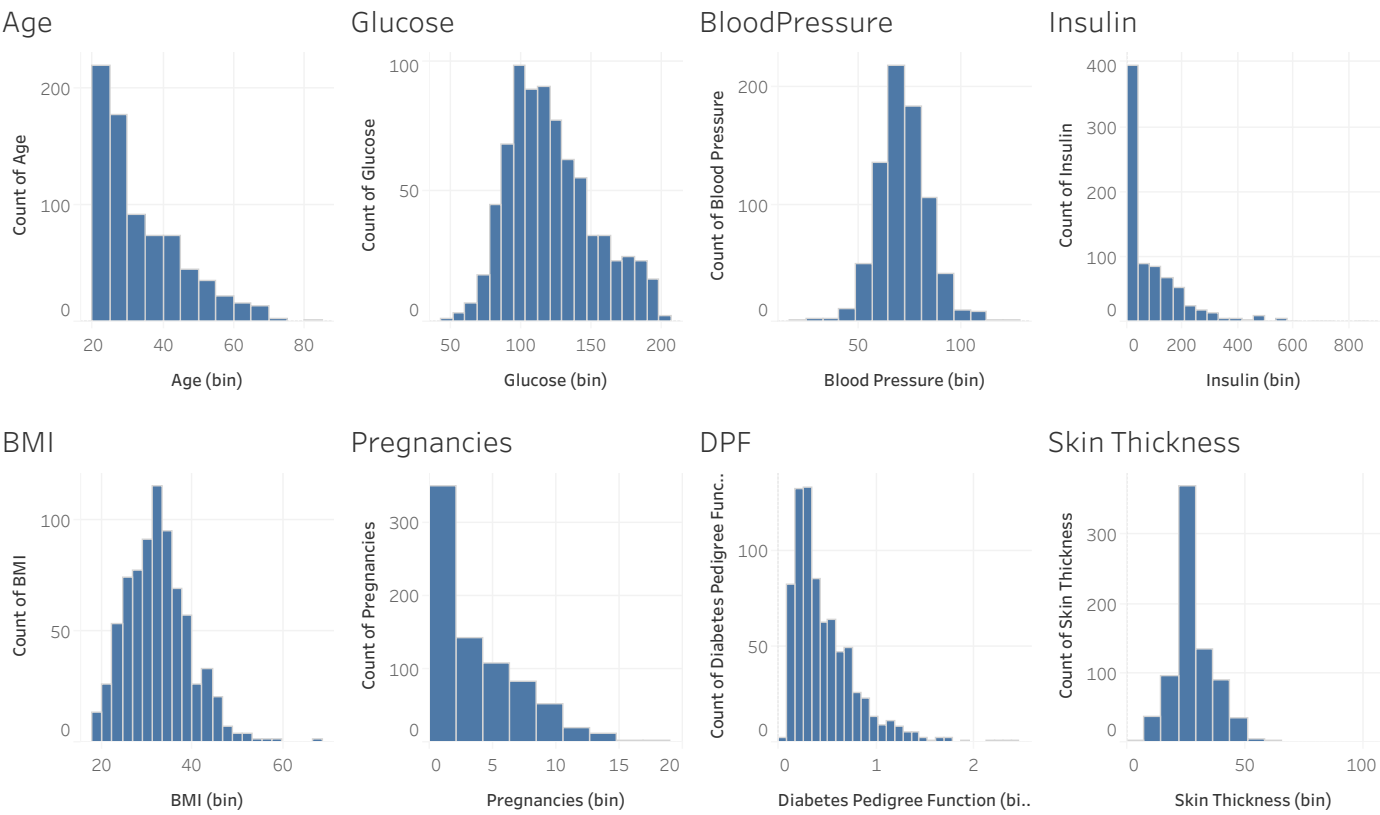
Analysis of Diabetic Patients

Pie chart to describe the diabetic or non-di..	Scatter charts between relevant var..	Histogram plot to see the distribution of th..	Operation on age column by creating th..
--	---------------------------------------	--	--



Analysis of Diabetic Patients


Pie chart to describe the diabetic or non-di..	Scatter charts between relevant var..	Histogram plot to see the distribution of th..	Operation on age column by creating th..
--	---------------------------------------	--	--




Analysis of Diabetic Patients

Pie chart to describe the diabetic or non-di..	Scatter charts between relevant var..	Histogram plot to see the distribution of th..	Operation on age column by creating th..
--	---------------------------------------	--	--

Avg age of Diabetic Patient



avg DPF for age group



Pregnancies with age groups

